

# Performance Modeling of Computer Systems and Networks

*Prof. Vittoria de Nitto Personè*

## Introduction to modeling

Università degli studi di Roma Tor Vergata  
Department of Civil Engineering and Computer Science Engineering

Copyright © Vittoria de Nitto Personè, 2021  
<https://creativecommons.org/licenses/by-nc-nd/4.0/>



1

### *Queueing theory*

is an area of mathematics, involving stochastic analysis, which allows one to **predict** the performance of a computer system and to **improve** performance

**Idea:** to analytically model the computer system as consisting of *resources* (like CPU, bandwidth, energy, VM, disk, etc.) and jobs which require these resources. Contention occurs when several jobs simultaneously require a resource, which means some jobs must wait.

Queueing theory allows you to predict what those “waits” will be and how to reduce it

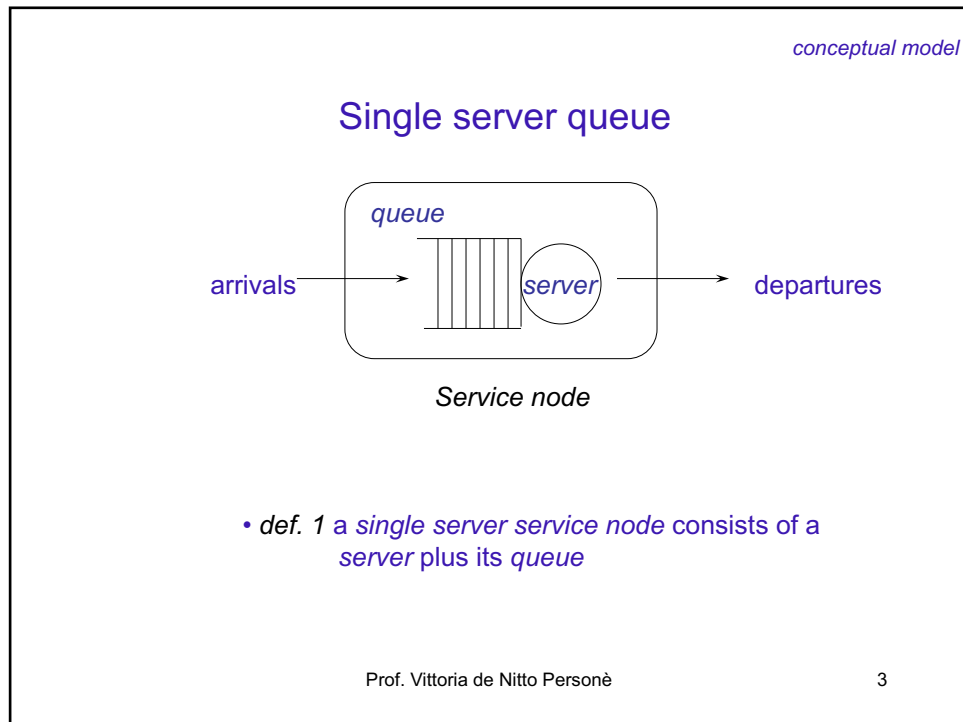


*So improving system performance!*

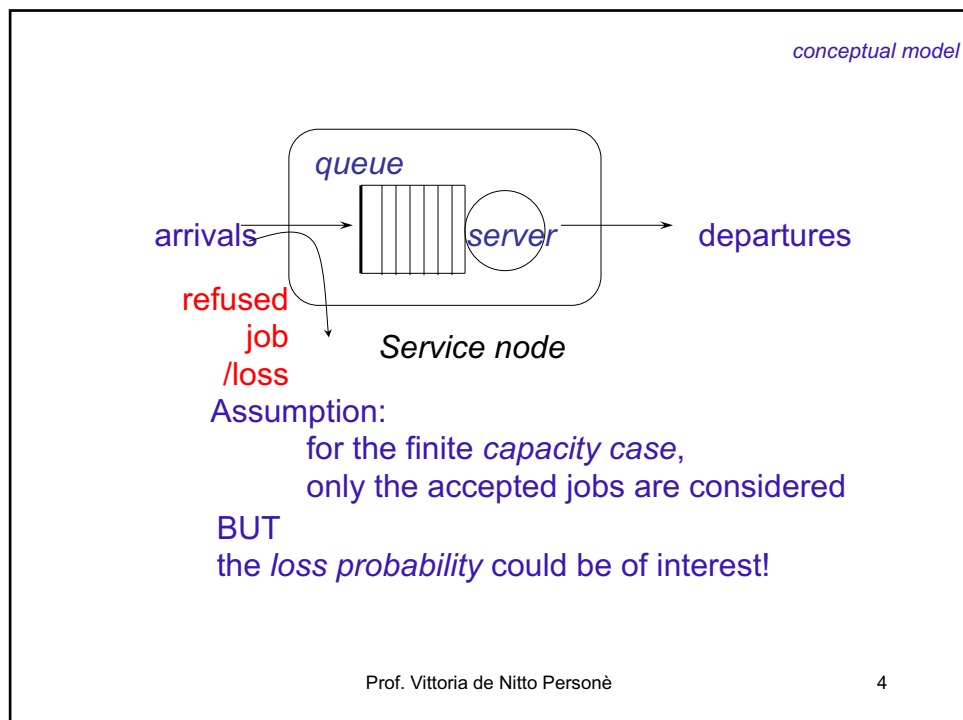
Prof. Vittoria de Nitto Personè

2

2



3



4

*def. 2 queue discipline (scheduling / service order):*  
the algorithm used when a job is selected  
from the queue to enter service

FIFO – first in, first out  
LIFO – last in, first out  
random – serve in random order  
Priority – typically shortest job first (SJF)  
PS – processor sharing

Prof. Vittoria de Nitto Personè

5

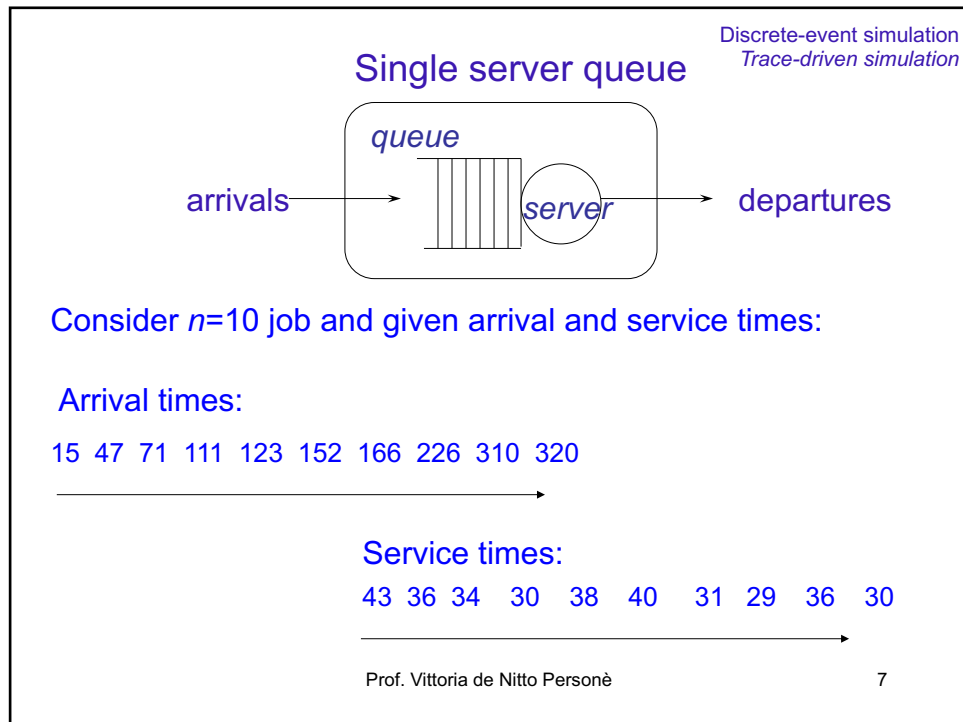
5

- FIFO (/ FCFS):
  - The order of arrival and departure are the same
  - A job cannot start service if the “previous” job has not left the node; this observation can be used to simplify the simulation
  - Unless otherwise specified, assume FIFO with infinite queue capacity
- service is *non-preemptive*
  - Once initiated, service of a job will continue until completion
- service is *conservative*
  - server will never remain idle if there is one or more jobs in the service node

Prof. Vittoria de Nitto Personè

6

6



7

Service times:

→

	mean	variance	resp time
43 36 34 30 38 40 31 29 36 30	34,7	20,21	26,70
63 16 54 10 18 60 51 9 56 10	34,7	504,21	32,70
9 10 10 16 18 51 54 56 60 63	=	=	12,80
63 60 56 54 51 18 16 10 10 9	=	=	77,70

Prof. Vittoria de Nitto Personè


8

8

### Modelling power: Predictor tool

consider that we're given

- a network
- some particular packet routes
- the packet arrival rates
- the transmission times
- wire lengths



*by queueing theory*

- the mean time packets spend waiting at a particular router  $i$ ,
- the distribution on the queue buildup at router  $i$
- the mean overall time to get from point A to point B in the network

Prof. Vittoria de Nitto Personè

9

9

### Modelling power: Design tool

system design is often a counter-intuitive process

Example 1: doubling arrival rate

Requests arrival



CPU  
Serve in  
FIFO order

*evolution of the configuration  
and workload (upgrade)*

- starting tomorrow the arrival rate will double.
- you should do whatever it takes to ensure that jobs experience the same mean response time. I.e., customers should not notice the effect of the increased arrival rate.

*Question:* By how much should you increase the CPU speed in order to maintain the same mean response time?

*Answer:* Less than double!


Prof. Vittoria de Nitto Personè

10

10

## Modelling power: Design tool

### Design Example 1: doubling arrival rate



doubling CPU speed together with doubling the arrival rate will result in cutting **the mean response time in half!**


Prof. Vittoria de Nitto Personè

11

11

## Modelling power: Design tool

### Design Example 1: doubling arrival rate



This is actually identical to the original system, but where time is sped up by a factor of 2 (i.e. a second of service in the original system now requires only half a second. Also, whereas in the old system 3 jobs per second arrive, now 6 jobs per second arrive)

*A faster time clock!*

**the mean response time becomes half**


Prof. Vittoria de Nitto Personè

12

12

## Modelling power: Design tool

Design Example 1: doubling arrival rate



Does the answer change?

*No!*

Prof. Vittoria de Nitto Personè

13

13


## Modelling power: Design tool

Design Example 2

*Resource*



Server



*Queue*


Prof. Vittoria de Nitto Personè

14

14

## Modelling power: Design tool

### Design Example 2



- The service time distribution is irrelevant


Prof. Vittoria de Nitto Personè

15

15

## Modelling power: Design tool

### Design Example 2



- Is the response time improved?
  - Is the throughput improved?
- Not really!

sometimes “improvements” do nothing

Prof. Vittoria de Nitto Personè


16

16



## Modelling power: Design tool

Design Example 2: sometimes “improvements” do nothing



*bottleneck identification*

- Is the response time improved?
- Is the throughput improved?

The already negligible effect goes to zero as  $N$  increases!


Prof. Vittoria de Nitto Personè

17

17

## Modelling power: Design tool

Design Example 2: sometimes “improvements” do nothing



- Is the response time improved?
- Is the throughput improved?

It is possible!  
If  $N$  is sufficiently low, then the “improvement” helps.  
Consider, for example, the case  $N = 1$ .


Prof. Vittoria de Nitto Personè

18

18

## Modelling power: Design tool

Design Example 2: sometimes “improvements” do nothing



- Is the response time improved?
- Is the throughput improved?

**Absolutely!**

now arrival times are independent of service completions


Prof. Vittoria de Nitto Personè

19

19

## Modelling power: Design tool

now arrival times  
are independent  
of service  
completions




Prof. Vittoria de Nitto Personè

20

20

## Modelling power: Design tool

Design Example 3: case study



Prof. Vittoria de Nitto Personè


21

21

## Modelling power: Design tool

Design Example 3: case study

During the observation interval  
( $T=900\text{sec}$ )




Prof. Vittoria de Nitto Personè

22

22

## Modelling power: Design tool

Design Example 3: case study



- CPU twice fast
- load balance between disks
- a second fast disk ?

The answer is *counter-intuitive!*

### Operational Analysis

- Easy
- Independent on distributional assumptions
- Do not need to know the full network topology

Which of the following changes is most effective in increasing throughput?


Prof. Vittoria de Nitto Personè

23

23

## Modelling power: Design tool

Design Example 4: One machine or many?



Goal: minimizing mean response time

Assumption: jobs *non-preemptible*  
each job must be run to completion

*hint: "it depends on the workload."*

depends on the variability of the job size distribution

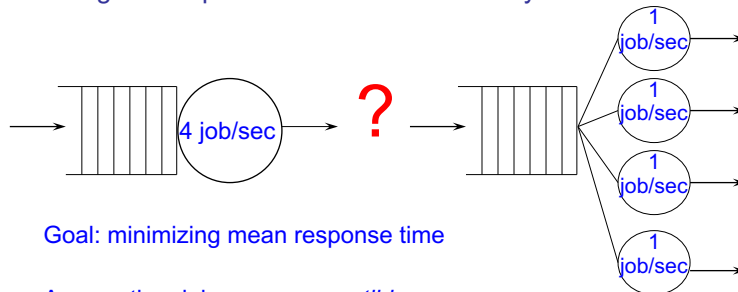
Prof. Vittoria de Nitto Personè

24

24

## Modelling power: Design tool

Design Example 4: One machine or many?



Goal: minimizing mean response time

Assumption: jobs *non-preemptible*  
each job must be run to completion

high variability  

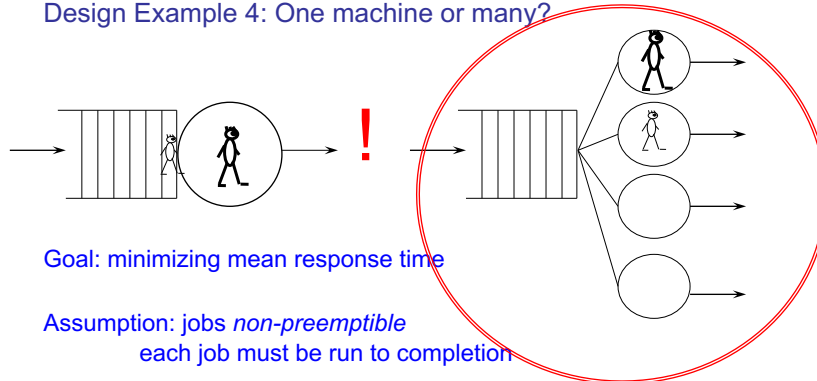
Prof. Vittoria de Nitto Personè

25

25

## Modelling power: Design tool

Design Example 4: One machine or many?



Goal: minimizing mean response time

Assumption: jobs *non-preemptible*  
each job must be run to completion

high variability  

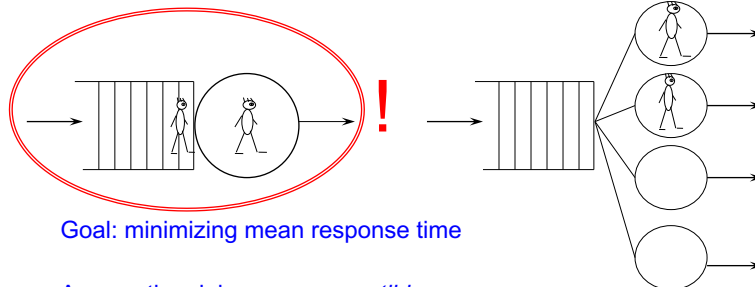
Prof. Vittoria de Nitto Personè

26

26

## Modelling power: Design tool

Design Example 4: One machine or many?



Goal: minimizing mean response time

Assumption: jobs *non-preemptible*  
each job must be run to completion

low variability



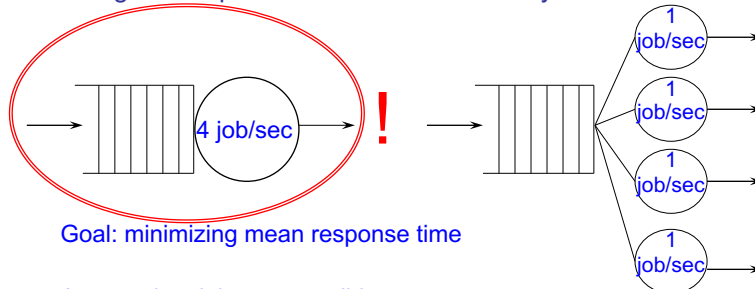
Prof. Vittoria de Nitto Personè

27

27

## Modelling power: Design tool

Design Example 4: One machine or many?



Goal: minimizing mean response time

Assumption: jobs *preemptible*  
each job can be stopped and restarted where they left off

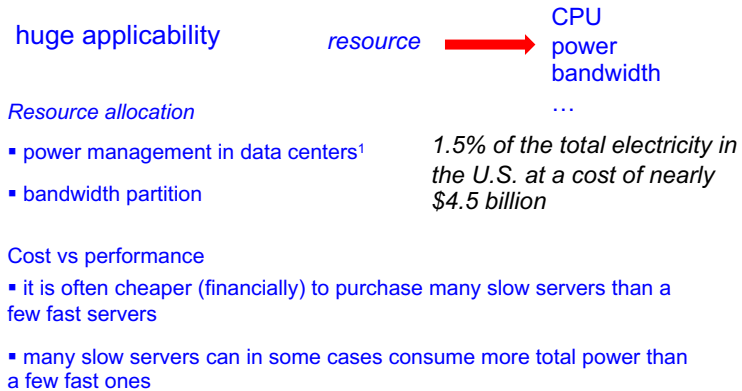
Prof. Vittoria de Nitto Personè

28

28

## Modelling power: Design tool

Design Example 4: One machine or many?



<sup>1</sup>A. Gandhi, M. Harchol-Balter, R. Das, and C. Lefurgy. Optimal power allocation in server farms. In *ACM Sigmetrics 2009 Conference on Measurement and Modeling of Computer Systems*, 2009.

Prof. Vittoria de Nitto Personè

29

29

*Dispositivi connessi: 1,2 miliardi 2018  
4,4 miliardi 2023*

*Consumo di energia dell'universo digitale, emissioni di CO2:  
2% 2008 → 3,7% 2020 → 8,5% 2025 → 14% 2040*

<https://www.corriere.it/dataroom-milena-gabanelli/emissioni-co2-ambiente-internet-quanto-inquina-nostra-vita-digitale-effetto-serra-consumi-invisibili-streaming-app-video/eb680526-5363-11eb-b612-933264f5acaf-va.shtml>

[https://www.facebook.com/watch/live/?v=148546343701326&ref=watch\\_permalink](https://www.facebook.com/watch/live/?v=148546343701326&ref=watch_permalink)


Prof. Vittoria de Nitto Personè

30

30

## Modelling power: Design tool

### Design Example 5: Task Assignment in a Server Farm




Prof. Vittoria de Nitto Personè

31

31

## Modelling power: Design tool

### Design Example 5: Task Assignment in a Server Farm



Objectives:

- web site "load-balancing"
- task assignment in distributed supercomputing servers

Prof. Vittoria de Nitto Personè


32

32



## Modelling power: Design tool

### Design Example 5: Task Assignment in a Server Farm



Assumption: homogeneous  
single resource for each job  
FIFO non-preemptible

Prof. Vittoria de Nitto Personè

33

33

## Modelling power: Design tool

### Design Example 5: Task Assignment in a Server Farm

*task assignment policies*

**Random** Each job flips a fair coin to determine where it is routed.


**Round-Robin** The  $i$ th job goes to host  $i \bmod n$ , where  $n$  is the number of hosts, hosts are numbered  $0, 1, \dots, n-1$ .

**Shortest-Queue** Each job goes to the host with the fewest number of jobs.

**Central-Queue** Rather than have a queue at each host, jobs are pooled at one central queue. When a host is done working on a job, it grabs the first job in the central queue to work on.

**Size-Interval-Task-Assignment (SITA)** "short" jobs go to the first host, "medium" jobs go to the second host, "long" jobs go to the third host, etc., for some definition of "short," "medium," "long." **job-size based**

**Least-Work-Left (LWL)** Each job goes to the host with the least total remaining work, where the "work" at a host is the sum of the sizes of jobs there.



Prof. Vittoria de Nitto Personè

34


34

## Modelling power: Design tool

### Design Example 5: Task Assignment in a Server Farm

#### task assignment policies

Which policies yields the lowest mean response time?



low variability → **LWL**

high variability → **SITA**

how important was it that we know the size of jobs?

Actually, most task assignment policies do not require knowing the size of jobs.  
(It can be proven by induction that LWL is equivalent to Central- Queue.)

Prof. Vittoria de Nitto Personè


35

35

## Modelling power: Design tool

### Design Example 5: Task Assignment in a Server Farm

#### task assignment policies



non-preemptible (supercomputer) → No-load balancing  
Different TAPs → different performance  
Well analyzed

preemptible (web server farm) → load balancing to minimize response time  
no-load balancing to minimize mean slowdown  
Open issue  
variable size distribution  
minimizing the variance of response time


Prof. Vittoria de Nitto Personè

36

36

## Modelling power: Design tool

### Design Example 6: Scheduling policies



- no-assumption on job size distribution
- non-preemptive jobs

Which of these will result in the lowest mean response time?

**First-In-First-Out (FIFO)** When the server completes a job, it starts working on the job which arrived earliest.

**Non-preemptive Last-In-First-Out (LIFO)** When the server completes a job, it starts working on job which arrived last.

**Random** When the server completes a job, it starts working on a random job.

These all have the same mean response time !!


Prof. Vittoria de Nitto Personè

37

37

## Modelling power: Design tool

### Design Example 6: Scheduling policies



~~**Non-preemptive Last-In-First-Out (LIFO)**~~ ~~When the server completes a job, it starts working on job which arrived last.~~

Whenever a new arrival enters the system, it immediately preempts the job in service

at least moderately variable  $\longrightarrow$  A huge performance improvement

hardly variable  $\longrightarrow$  Up to a factor of 2 worsening

Prof. Vittoria de Nitto Personè

38

38