

Performance Modeling of Computer Systems and Networks

Prof. Vittoria de Nitto Personè

Case study 1 *A single server queue*

Università degli studi di Roma Tor Vergata
Department of Civil Engineering and Computer Science Engineering


Copyright © Vittoria de Nitto Personè, 2021
<https://creativecommons.org/licenses/by-nc-nd/4.0/>



1

Single server queue

Discrete-event simulation
Trace-driven simulation



Consider $n=10$ job and given arrival and service times:

Arrival times:

15 47 71 111 123 152 166 226 310 320



Service times:

43 36 34 30 38 40 31 29 36 30



Prof. Vittoria de Nitto Personè

2

2

Discrete-event simulation
Trace-driven simulation

Algorithm 2: *trace-driven simulation*

```

c0 = 0.0;                                /* assumes that a0 = 0.0 */
i = 0;
while ( more jobs to process ) {
    i++;
    ai = GetArrival();
    if (ai < ci-1) di = ci-1 - ai;
    else di = 0.0;
    si = GetService();
    ci = ai + di + si;
}
n = i;
return d1, d2, . . . , dn;

```

Read data from a file

Prof. Vittoria de Nitto Personè

3

3

ssq1.c

```

#include <stdio.h>

#define FILENAME "ssq1.dat" /* input data file */
#define START 0.0

double GetArrival(FILE *fp) /* read an arrival time */
{
    double a;
    fscanf(fp, "%lf", &a);
    return (a);}

double GetService(FILE *fp) /* read a service time */
{
    double s;
    fscanf(fp, "%lf\n", &s);
    return (s);}

```

4

```

int main(void)
{ FILE *fp; /* input data file */
  long index = 0; /* job index */
  double arrival = START; /* arrival time*/
  double delay; /* delay in queue*/
  double service; /* service time*/
  double wait; /* delay + service*/
  double departure = START; /* departure time*/
  struct { /* sum of ... */
    double delay; /*delay times */
    double wait; /*wait times*/
    double service; /*service times */
    double interarrival; /* interarrival times */
  } sum = {0.0, 0.0, 0.0};

```

Prof. Vittoria de Nitto Personè

5

5

```

fp = fopen(FILENAME, "r");
if (fp == NULL) {
  fprintf(stderr, "Cannot open input file %s\n", FILENAME);
  return (1); }
while (!feof(fp)) {
  index++;
  arrival = GetArrival(fp);
  if (arrival < departure)
    delay = departure - arrival; /* delay in queue */
  else delay = 0.0; /* no delay */
  service = GetService(fp);
  wait = delay + service;
  departure = arrival + wait; /* time of departure */
  sum.delay += delay;
  sum.wait += wait;
  sum.service += service; }

```

Prof. Vittoria de Nitto Personè

6

6

```

sum.interarrival = arrival - START;

printf("\nfor %ld jobs\n", index);
printf("average interarrival time = %.2f\n", sum.interarrival / index);
printf("    average service time .... = %.2f\n", sum.service / index);
printf("    average delay ..... = %.2f\n", sum.delay / index);
printf("    average wait ..... = %.2f\n", sum.wait / index);

fclose(fp);
return (0);}

```

Prof. Vittoria de Nitto Personè

7

7

Discrete-event simulation
Output statistics

		Output state									
	i	1	2	3	4	5	6	7	8	9	10
read from file	a_i	15	47	71	111	123	152	166	226	310	320
from algorithm	d_i	0	11	23	17	35	44	70	41	0	26
read from file	s_i	43	36	34	30	38	40	31	29	36	30

- average interarrival time $\bar{r} = \frac{a_n}{n} = \frac{320}{10} = 32.0$ sec
- average service time $\bar{s} = 34.7$ sec
- arrival rate $\frac{1}{\bar{r}} \approx 0.031$ job/sec
- service rate $\frac{1}{\bar{s}} \approx 0.029$ job/sec
- traffic intensity $\frac{\bar{s}}{\bar{r}} = 1.084375$
- utilization $\bar{x} = \frac{n}{c_n} \bar{s} = 0.92287$

Insight: The server is not quite able to process jobs at the rate they arrive on average.

Prof. Vittoria de Nitto Personè

8

8

Discrete-event simulation
Output statistics

	i	1	2	3	4	5	6	7	8	9	10
read from file	a_i	15	47	71	111	123	152	166	226	310	320
from algorithm	d_i	0	11	23	17	35	44	70	41	0	26
read from file	s_i	43	36	34	30	38	40	31	29	36	30
completion time		58	94	128	158	196	236	267	296	346	376

Prof. Vittoria de Nitto Personè

9

Discrete-event simulation
Output statistics

Job-averaged statistics

- average delay and service time $\bar{d} = 26.7$, $\bar{s} = 34.7$ sec
therefore, the average wait time is:

$$\bar{w} = \bar{d} + \bar{s} = 26.7 + 34.7 = 61.4 \text{ sec}$$
- verification is a difficult step
 ↓
 Consistency check:
 used to verify that a simulation satisfies known equations
 - compute $\bar{w}, \bar{d}, \bar{s}$ *independently*
 - then verify that $\bar{w} = \bar{d} + \bar{s}$

Prof. Vittoria de Nitto Personè

10

$$\bar{l} = \frac{1}{\tau} \int_0^{\tau} l(t) dt \quad \bar{q} = \frac{1}{\tau} \int_0^{\tau} q(t) dt \quad \bar{x} = \frac{1}{\tau} \int_0^{\tau} x(t) dt$$

With $\tau = c_{10} = 376$

$$\bar{l} = \frac{n}{c_n} \bar{w} \quad \bar{l} = \frac{10}{376} 61.4 = 1.633 \quad \bar{q} = 0.710 \quad \bar{x} = 0.923$$

The average of numerous random observations (samples) of the number in the service node should be close to \bar{l} .
(Same holds for \bar{q}, \bar{x})

11

Service times:

43 36 34 30 38 40 31 29 36 30



63 16 54 10 18 60 51 9 56 10 *alta variabilità*

9 10 10 16 18 51 54 56 60 63 *fair*

63 60 56 54 51 18 16 10 10 9 *unfair*

12

Dati originali

for 10 jobs

average interarrival time = 32.00
 average service time = 34.70
 average delay = 26.70
 average wait = 61.40

$$\bar{l} = \frac{10}{376} 61.4 = 1.633$$

Ordinamento unfair

for 10 jobs

average interarrival time = 32.00
 average service time = 34.70
 average delay = 77.70
 average wait = 112.40

$$\bar{l} = \frac{10}{376} 112.4 = 2.989$$

Dati con alta variabilità

for 10 jobs

average interarrival time = 32.00
 average service time = 34.70
 average delay = 32.70
 average wait = 67.40

$$\bar{l} = \frac{10}{376} 67.4 = 1.793$$

Ordinamento fair

for 10 jobs


average interarrival time = 32.00
 average service time = 34.70
 average delay = 12.80
 average wait = 47.50

$$\bar{l} = \frac{10}{376} 47.50 = 1.263$$

Prof. Vittoria de Nitto Personè

13

13



Prof. Vittoria de Nitto Personè

14

14

Case Study

- The owners of an ice cream shop are considering adding additional flavors and cone options
- But they are concerned about the effect of the resultant increase in service times on queue length

The case can be studied as a single-server queue

- ssq1.dat represents 1000 customer interactions at the shop (arrival times and the corresponding service times)
- for the study, the service times are systematically increased (and decreased) by a common multiplicative factor

Exercise 1

By running ssq1 for the datafile ssq1.dat, the following can be observed:

Observed arrival rate $\frac{1}{\bar{r}} \approx 0.10$ job/sec

Observed service rate $\frac{1}{\bar{s}} \approx 0.14$ job/sec

- Modify the program to compute $\bar{l}, \bar{q}, \bar{x}$

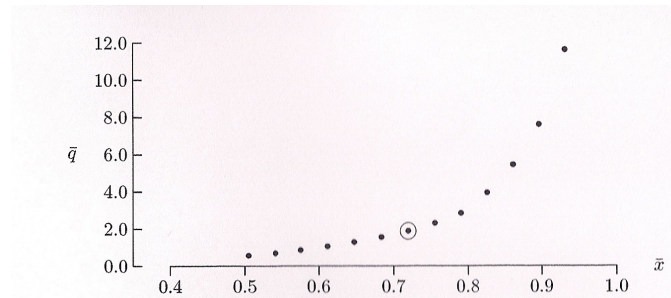
- You will find $1 - \bar{x} \approx 0.28$

Despite this significant idle time (28%), enough jobs are delayed so that the average number in the queue is

$$\bar{q} \approx 2$$

Discrete-event simulation
ssq - case study

15% ↑ service times 109% ↑ queue length
30% ↑ service times 518% ↑ queue length




- even a modest increase in service times will produce a significant increase in the average queue length
- nonlinear relationship between \bar{q}, \bar{x} particularly pronounced when utilization approaches 1

Prof. Vittoria de Nitto Personè

17

17

Discrete-event simulation
ssq - case study



- \bar{q}, \bar{x} are continuous in the values → many additional points more
- few would question the validity of “connecting the dots”

Prof. Vittoria de Nitto Personè

18

18

Guidelines

- If the data have essentially no uncertainty and the resulting is smooth
OK connecting the dots, but the original are left
- If the data have essentially no uncertainty but the resulting is not
smooth more dots need to be generated
- If the dots correspond to uncertain data, NO interpolation!
(approximation)
- If the data is inherently discrete, NEVER interpolation!

Exercises

- Ex 1.1.2 and Ex 1.1.3 on p.11 from textbook
- Study the program ssq1.c
- Run with ssq1.dat; analyze the results
- Generate a data file with the values on p.8
- Run with the new data, verify that the results confirm the expectations
- Ex 1.2.2 and Ex 1.2.3 on p.24 of the textbook