



Performance Modeling of Computer Systems and Networks

Prof. Vittoria de Nitto Personè

Discrete-Event Simulation
examples

Università degli studi di Roma Tor Vergata
Department of Civil Engineering and Computer Science Engineering


Copyright © Vittoria de Nitto Personè, 2021
<https://creativecommons.org/licenses/by-nc-nd/4.0/>



1

Discrete-Event Simulation
case study ssq

Single Server Queue



Arrival times: a_i

~~15 47 71 111 123 152 166 226 310 320~~

Pseudo-random generators

Service times: s_i


~~43 36 34 30 38 40 31 29 36 30~~

Prof. Vittoria de Nitto Personè

2

2

Single Server Queue



- assume **service times** are between 1.0 and 2.0 minutes
 - The distribution within this range is unknown
 - Without further knowledge, we assume no time is more likely than any other
- Uniform(1.0, 2.0)

Prof. Vittoria de Nitto Personè


3

3

Exponential distribution

- In general, it is unreasonable to assume that all possible values are equally likely
- Frequently, small values are more likely than large values
- We need a non-linear transformation that maps $0.0 \rightarrow 1.0$ to $0.0 \rightarrow \infty$

this is the most frequently used function
 $\mu > 0$ is a parameter that "control" the frequency of large values in respect of the small ones



Prof. Vittoria de Nitto Personè

4

4


- the transformation is monotone increasing, one-to-one

$$\begin{aligned}
 0 < u < 1 &\Leftrightarrow 0 < (1 - u) < 1 \\
 &\Leftrightarrow -\infty < \ln(1 - u) < 0 \\
 &\Leftrightarrow 0 < -\mu \ln(1 - u) < \infty \\
 &\Leftrightarrow 0 < x < \infty
 \end{aligned}$$

```
double Exponential(double μ)    /* use μ > 0.0 */
{
    return (-μ * log(1.0 - Random()));
}
```

- the parameter μ specifies the sample mean

Single Server Queue



Arrival times: a_i

- use the exponential function for the interarrival times

$$a_i = a_{i-1} + \text{Exponential}(\mu); i = 1, 2, 3, \dots, n$$

Service times: s_i

Uniform(1.0, 2.0)

- program ssq2 is an extension of ssq1
 - arrival times are drawn from *Exponential*(2.0)
 - service times are drawn from *Uniform*(1.0, 2.0)

trace-driven simulation

```
#include <stdio.h>

#define FILENAME  "ssq1.dat"    /* input data file */
#define START     0.0

double GetArrival(FILE *fp)    /* read an arrival time */
{
    double a;
    fscanf(fp, "%lf", &a);
    return (a);}

double GetService(FILE *fp)    /* read a service time */
{
    double s;
    fscanf(fp, "%lf\n", &s);
    return (s);}
```

ssq2.c distribution-driven simulation

```

#include <stdio.h>
#include <math.h>
#include "rng.h"
#define LAST      10000L    /* number of jobs processed */
#define START     0.0

double Exponential(double m)           /* -----*
{return (-m * log(1.0 - Random())); }    m > 0.0
                                         -----*/

double Uniform(double a, double b)     /* -----*
{return (a + (b - a) * Random()); }      a < b
                                         * -----*/

double GetArrival(void)
{static double arrival = START;
 arrival += Exponential(2.0);
 return (arrival);}

double GetService(void)
{return (Uniform(1.0, 2.0));}

```

Prof. Vittoria de Nitto Personè

9

9

Discrete-Event Simulation case study ssq

- the program generates all first-order statistics

$$\bar{r}, \bar{w}, \bar{d}, \bar{s}, \bar{l}, \bar{q}, \bar{x}$$

Prof. Vittoria de Nitto Personè

10

10

trace-driven simulation = distribution-driven simulation

```
int main(void)
{ FILE *fp; /* input data file */
  long index = 0; /* job index */
  double arrival = START; /* arrival time*/
  double delay; /* delay in queue*/
  double service; /* service time*/
  double wait; /* delay + service*/
  double departure = START; /* departure time*/
  struct { /* sum of... */
    double delay; /*delay times */
    double wait; /*wait times*/
    double service; /*service times */
    double interarrival; /* interarrival times */
  } sum = {0.0, 0.0, 0.0};
```

Prof. Vittoria de Nitto Personè

11

11

trace-driven simulation

```
fp = fopen(FILENAME, "r");
if (fp == NULL) {
  fprintf(stderr, "Cannot open input file %s\n", FILENAME);
  return (1); }
while (!feof(fp)) {
```

distribution-driven simulation

```
PutSeed(123456789);
while (index < LAST) {
```

Prof. Vittoria de Nitto Personè

12

12

```

                                trace-driven: fp
while (index < LAST) {
    index++;
    arrival      = GetArrival();
    if (arrival < departure)
        delay = departure - arrival; /* delay in queue */
    else delay = 0.0;                /* no delay */
    service = GetService();
    wait = delay + service;
    departure = arrival + wait; /* time of departure */
    sum.delay += delay;
    sum.wait += wait;
    sum.service += service;
}
sum.interarrival = arrival - START;

```

Prof. Vittoria de Nitto Personè

13

13

```

printf("\nfor %ld jobs\n", index);
printf("  average interarrival time = %6.2f\n", sum.interarrival / index);
printf("  average wait ..... = %6.2f\n", sum.wait / index);
printf("  average delay ..... = %6.2f\n", sum.delay / index);
printf("  average service time .... = %6.2f\n", sum.service / index);
printf("  average # in the node ... = %6.2f\n", sum.wait / departure);
printf("  average # in the queue .. = %6.2f\n", sum.delay / departure);
printf("  utilization ..... = %6.2f\n", sum.service / departure);
return (0);

```

Prof. Vittoria de Nitto Personè

14

14

Example 1

- The “theoretical” averages using *Exponential*(2.0) (rate 0.5 j/s) arrivals and *Uniform*(1.0, 2.0) (rate 0.67) service times are

\bar{r}	\bar{w}	\bar{d}	\bar{s}	\bar{l}	\bar{q}	\bar{x}	
2.00	3.83	2.33	1.50	1.92	1.17	0.75	exact analytical results, No simulation!


- Although the server is busy 75% of the time, on average there are approximately 2 jobs in the service node
- A job can expect to spend more time in the queue than in service
- To achieve these averages, many jobs must pass through node

Prof. Vittoria de Nitto Personè

15

15

- The accumulated average wait was printed every 20 job



- The convergence of \bar{w} is slow, erratic, and dependent on the initial seed

Prof. Vittoria de Nitto Personè

16

16

- the program can be used to study the **steady-state** behavior
 - Will the statistics converge independent of the initial seed?
 - How many jobs does it take to achieve steady-state behavior?
- the program can be used to study the **transient** behavior
 - Fix the number of jobs processed and replicate the program with the initial state fixed
 - Each replication uses a different initial rng seed

17

Steady-state analysis

	\bar{r}	\bar{w}	\bar{d}	\bar{s}	\bar{l}	\bar{q}	\bar{x}	
<u>theoretical</u>	2.00	3.83	2.33	1.50	1.92	1.17	0.75	
n=10	2.85	1.74	0.39	1.35	0.59	0.13	0.45	seed=123456789
n=100	2.06	3.16	1.67	1.48	1.50	0.80	0.71	
n=1000	2.03	3.44	1.94	1.50	1.69	0.96	0.74	
n=10000	2.02	3.86	2.36	1.50	1.91	1.17	0.74	
n=100000	2.00	3.85	2.35	1.50	1.92	1.17	0.75	
n=1000000	2.00	3.81	2.31	1.50	1.90	1.15	0.75	
n=1000000	2.00	3.84	2.34	1.50	1.92	1.17	0.75	seed=1

n=10	2.13	2.36	0.75	1.62	1.02	0.32	0.69	seed=1
n=10	1.48	2.37	0.87	1.50	1.24	0.46	0.79	seed=987654321
n=10	1.49	1.89	0.49	1.40	1.12	0.29	0.83	seed=21212121

Transient analysis

18

Simulating an initial steady state

$$\bar{d} = 2.33$$

$$\text{departure}=3$$

$$a_1 = a_0 + \text{expo}(2) = 0 + 0.8 = 0.8$$

$$0.8 < 3$$

$$d_1 = 3 - 0.8 = 2.2$$

$$s_1 = \text{Uniform}(1,2) = 1.3$$

$$w_1 = 2.2 + 1.3 = 3.5$$

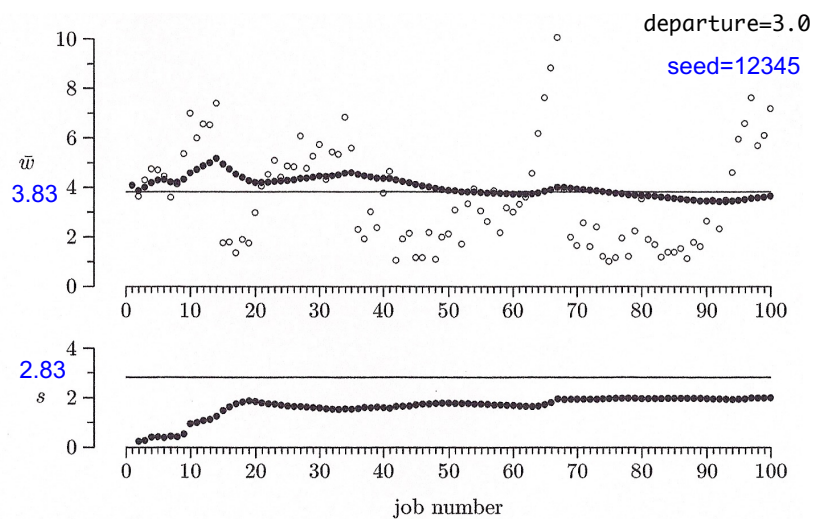
$$c_1 = 0.8 + 3.5 = 4.3$$

Prof. Vittoria de Nitto Personè

19

19


DE simulation
Sample statistics



Prof. Vittoria de Nitto Personè

20

20



21

DE simulation
Sample statistics

Serial correlation

- *Independence*: each x_i value does not depend on any other point
- Time-sequenced DES output is typically not independent
- E.g.: wait times of consecutive jobs have positive *serial correlation*
- Example: Consider output from ssq2
 - *Exponential(2)* interarrivals, *Uniform(1,2)* service
- wait times w_1, w_2, \dots, w_{100} , have high positive serial correlation
 - The correlation produces a bias in the standard deviation

Prof. Vittoria de Nitto Personè

22

22

Example 2

- assume that jobs arrive at random with a steady-state arrival rate of 0.5 jobs per minute
- assume that Job service times are “composite” with two components:
 - the *number* of service tasks is $1 + \text{Geometric}(0.9)$
 - the *time* (in minutes) per task is $\text{Uniform}(0.1, 0.2)$

```
double GetService(void)
{
    long k;
    double sum = 0.0;
    long tasks = 1 + Geometric(0.9);
    for (k = 0; k < tasks; k++)
        sum += Uniform(0.1, 0.2);
    return (sum);
}
```

Prof. Vittoria de Nitto Personè

23

23

- The theoretical steady-state statistics for this model are

\bar{r}	\bar{w}	\bar{d}	\bar{s}	\bar{l}	\bar{q}	\bar{x}	
2.00	5.77	4.27	1.50	2.89	2.14	0.75	exact analytical results, No simulation!
	(3.83	2.33		1.92	1.17)		

- The arrival rate, service rate, and utilization are identical to the previous case (example 1)
- The other four statistics are significantly larger
- performance measures are **sensitive** to the choice of service time distribution

Prof. Vittoria de Nitto Personè

24

24

A simple inventory system



- Distributes items from current inventory to customers
- Customer demand is discrete
- Simple: one type of item
 - Inventory review is periodic
 - Items are ordered, if necessary, only at review times
 - (s, S) are the min,max inventory levels, $0 \leq s < S$

Prof. Vittoria de Nitto Personè

25

25

Simply Inventory System

- Program sis2 has randomly generated demands using an *Equilikely*(a, b) random variate
- Using random data, we can study transient and steady-state behaviors

#include <stdio.h>

sis1.c

~~#define FILENAME "sis1.dat"~~

#define MINIMUM 20

#define MAXIMUM 80

#define STOP 100

#define sqr(x) ((x) * (x))

~~long GetDemand(FILE *fp)~~

{

long d;

~~fscanf(fp, "%ld\n", &d);~~~~return (d);}~~

Prof. Vittoria de Nitto Personè

26

26

sis2.c

```

#include <stdio.h>
#include "rng.h"

#define MINIMUM 20
#define MAXIMUM 80
#define STOP 100 /* 100 weeks = about 2 years*/
#define sqr(x) ((x) * (x))

long Equilikely(long a, long b)
{ return (a + (long) ((b - a + 1) * Random()));}

long GetDemand(void)
{
    return (Equilikely(10, 50)); }


    i : 1 2 3 4 5 6 7 8 9 10 11 12
    di : 30 15 25 15 45 30 25 15 20 35 20 30

```

27

```

int main(void)
{
    long index = 0;
    long inventory = MAXIMUM;
    long demand;
    long order;
    struct {
        double setup;
        double holding; /*inventory hold (+) */
        double shortage; /*inventory short (-) */
        double order;
        double demand;
    } sum = { 0.0, 0.0, 0.0, 0.0, 0.0 };

    PutSeed(123456789); 

```

Prof. Vittoria de Nitto Personè

28

28

```

while (index < STOP) {
    index++;
    if (inventory < MINIMUM) {
        order = MAXIMUM - inventory;
        sum.setup++;
        sum.order += order;
    }
    else order = 0;
    inventory += order; /* there is no delivery lag */ demand =
    GetDemand();
    sum.demand += demand;
    if (inventory > demand)
        sum.holding += (inventory - 0.5 * demand);
    else {
        sum.holding += sqr(inventory) / (2.0 * demand);
        sum.shortage += sqr(demand - inventory) / (2.0 * demand);
    }
    inventory -= demand;
}

```

Prof. Vittoria de Nitto Personè

29

29

```

if (inventory < MAXIMUM) {
    order = MAXIMUM - inventory;
    sum.setup++;
    sum.order += order;
    inventory += order;
}

...

```

Prof. Vittoria de Nitto Personè

30

30

for 100 time intervals with an average demand of 27.68
and policy parameters (s, S) = (20, 80)

average order = 27.68
setup frequency = 0.36
average holding level = 44.81
average shortage level ... = 0.14

- trace file sis1.dat contains data for n=100 time intervals
- with (s, S)=(20, 80)

$$\bar{o} = \bar{d} = 29.29 \quad \bar{u} = 0.39 \quad \bar{l}^+ = 42.40 \quad \bar{l}^- = 0.25$$

Prof. Vittoria de Nitto Personè

31

31

```
int main(void)
{ long seed;
  long index      = 0;
  long inventory = MAXIMUM;
  long demand;
  long order;
  struct {
    double setup;
    double holding; /*inventory held (+) */
    double shortage; /*  inventory short (-)  */
    double order;
    double demand;
  } sum = { 0.0, 0.0, 0.0, 0.0, 0.0 };

  PutSeed(-1);
  GetSeed(&seed);
  printf("\nwith an initial seed of %ld", seed);
```

Prof. Vittoria de Nitto Personè

32

32

for 100 time intervals with an average demand of 27.68
and policy parameters $(s, S) = (20, 80)$

average order = 27.68
setup frequency = 0.36
average holding level = 44.81
average shortage level ... = 0.14

with an initial seed of 1333437895
for 100 time intervals with an average demand of 31.00
and policy parameters $(s, S) = (20, 80)$

average order = 31.00
setup frequency = 0.40
average holding level = 43.39
average shortage level ... = 0.37

Prof. Vittoria de Nitto Personè

33

33

Discrete-Event Simulation
case study sis

Simply Inventory System

if $(a, b) = (10, 50)$ and $(s, S) = (20, 80)$, then the approximate
steady-state statistics are

\bar{d}	$\bar{\sigma}$	\bar{u}	\bar{l}^+	\bar{l}^-
30.00	30.00	0.39	42.86	0.26

(trace-driven

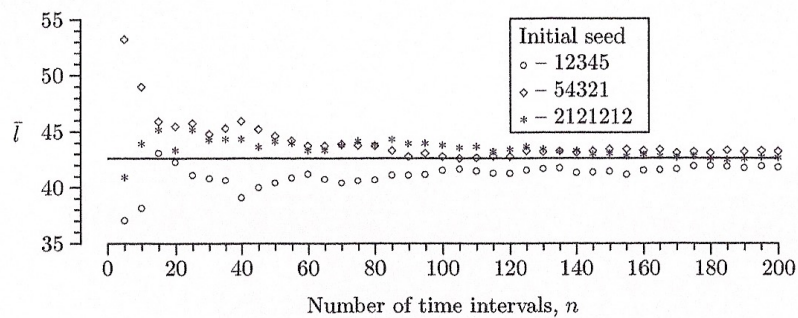
$\bar{\sigma} = \bar{d} = 29.29$ $\bar{u} = 0.39$ $\bar{l}^+ = 42.40$ $\bar{l}^- = 0.25$)

Prof. Vittoria de Nitto Personè

34

34

The average inventory level $\bar{I} = \bar{I}^+ - \bar{I}^-$
approaches steady state after several hundred time intervals



Convergence is slow, erratic, and dependent on the initial seed

Prof. Vittoria de Nitto Personè

35

35


- using a fixed initial seed guarantees the exact same demand sequence (in the example 12345)
- any changes to the system are caused solely by the change of s
- a steady state study of this system is unreasonable:
 - all parameters would have to remain fixed for many years
 - when $n=100$, we simulate approximately 2 years
 - when $n=10000$, we simulate approximately 192 years

Prof. Vittoria de Nitto Personè

36

36

- if we fix S , we can find the optimal cost by varying s
- recall that the dependent cost ignores the fixed cost of each item



Prof. Vittoria de Nitto Personè

37

37

Statistical Considerations (sect. 3.1.3)

- *variance reduction:*
(intuitive approach) use the same random numbers
(We kept the same initial seed and changed only s)
- NOTE:
transient behavior will always have some inherent uncertainty
- *Robust Estimation:*
when a data point is not sensitive to small changes in assumptions
 - values of s close to 23 have essentially the same cost
 - Would the cost be more sensitive to changes in S or other assumed values?

Prof. Vittoria de Nitto Personè

38

38

Exercises

- derive analytical results on p.12
- study program ssq2.c; run it and compare output with the results on p.12
- Exercises: 3.1.1, 3.1.2, 3.1.4, 3.1.5, 3.1.6