

Performance Modeling of Computer Systems and Networks

Prof. Vittoria de Nitto Personè

Analytical results
KP further results

Università degli studi di Roma Tor Vergata
Department of Civil Engineering and Computer Science Engineering


Copyright © Vittoria de Nitto Personè, 2021
<https://creativecommons.org/licenses/by-nc-nd/4.0/>



1

Analytical models
M/G/1

The waiting time and the remaining service time



$$T_Q = t_c \oplus S_{rem}$$


$$E(S_{rem}) = \frac{\lambda}{2} E(S^2)$$

Prof. Vittoria de Nitto Personè

2

2

The waiting time and the remaining service time



$$E(S_{rem}) = \frac{\lambda}{2} E(S^2)$$

exponential $\longrightarrow E(S^2) = 2E(S)^2 \longrightarrow E(S_{rem}) = \frac{\lambda}{2} 2E(S)^2$


$$E(S_{rem}) = \rho E(S)$$

Prof. Vittoria de Nitto Personè

3

3

The waiting time and the remaining service time



$$E(T_Q) = \frac{\rho E(S)}{1 - \rho} = \frac{E(S_{rem})}{1 - \rho} = \frac{1}{1 - \rho} E(S_{rem})$$

exponential $E(S_{rem}) = \rho E(S)$

Prof. Vittoria de Nitto Personè

4

4

$$\begin{aligned}
 E(T_Q) &= \frac{\rho}{1-\rho} \frac{C^2+1}{2} E(S) = \\
 &= \frac{\rho}{2(1-\rho)} \left[\frac{\sigma^2(S)}{E(S)^2} + 1 \right] E(S) = \\
 &= \frac{\rho}{2(1-\rho)} \left[\frac{E(S^2) - E(S)^2}{E(S)^2} + 1 \right] E(S) = \\
 &= \frac{\lambda E(S)}{2(1-\rho)} \left[\frac{E(S^2)}{E(S)^2} - 1 + 1 \right] E(S) = \\
 &= \frac{\lambda}{2(1-\rho)} \left[\frac{E(S^2)}{E(S)^2} \right] E(S)^2 = \frac{\frac{\lambda}{2} E(S^2)}{1-\rho}
 \end{aligned}$$


Prof. Vittoria de Nitto Personè

5

5

Analytical models
M/G/1

The waiting time and the remaining service time



$$T_Q = t_c \oplus S_{rem}$$

$$\frac{1}{1-\rho}$$

represents the mean time to complete the jobs in the queue at the arrival instant

$$E(S_{rem})$$


is the mean time to complete the job in service at the arrival instant

Prof. Vittoria de Nitto Personè

6

6

The waiting time and the remaining service time




Prof. Vittoria de Nitto Personè

7

7

The waiting time and the remaining service time




Prof. Vittoria de Nitto Personè

8

8

The waiting time and the remaining service time




$$E(T_Q) = \frac{E(S_{rem})}{1-\rho} = \frac{\frac{\lambda}{2}E(S^2)}{1-\rho} \quad \text{M/G/1}$$

Prof. Vittoria de Nitto Personè

9

9

The response time



$$E(T_S) = E(T_Q) + E(S) = \frac{\frac{\lambda}{2}E(S^2)}{1-\rho} + E(S) \quad \text{M/G/1}$$

$$E(T_S) = \frac{\rho E(S)}{1-\rho} + E(S) = \frac{E(S)}{1-\rho} \quad \text{M/M/1}$$

Prof. Vittoria de Nitto Personè

10

10

Non-preemptive abstract scheduling

Def. 1

A policy is *preemptive* if a job may be stopped part way through its execution and then resumed at a later point in time from the same point where it was stopped. A policy is *non-preemptive* if jobs are always run-to-completion.

Def. 2

A *work-conserving* scheduling policy is one which always performs work on some job when there is a job in the system.

Theorem 1 (Conway, Maxwell, Miller¹).

All non-preemptive service orders that do not make use of job sizes have the same distribution on the number of jobs in the system.

$$E(N_S) \quad E(T_S) \quad E(N_Q) \quad E(T_Q)$$


¹Richard Conway, William Maxwell, and Louis Miller, Theory of Scheduling Addison-Wesley Publishing Company, Inc., 1967. Chapter 8

Non-preemptive abstract scheduling

$$E(T_Q) = \frac{\frac{\lambda}{2} E(S^2)}{1 - \rho}$$

which is very high when $E(S^2)$ is high

WHAT IS FAIR?



13

Let us consider the mean time in system for a job of size x

$$E(T_S(x)) = E(x + T_Q(x)) = x + E(T_Q) = x + \frac{\frac{\lambda}{2} E(S^2)}{1 - \rho}$$

14

Analytical models
scheduling

WHAT IS FAIR?

Response time is proportional to size

Everyone waits the same amount

Define the slowdown for a job of size x as

$$E(sd(x)) = \frac{E(T_S(x))}{x}$$

15

Analytical models
scheduling

Slowdown for jobs of size x

Def.
The mean slowdown for jobs of size x is the observed mean response time in respect of their size, that is

$$E(sd(x)) = \frac{E(T_S(x))}{x}$$

$$E(sd(x)) = 1 + \frac{\frac{\lambda}{2} E(S^2)}{x(1-\rho)}$$

Note that small jobs have a higher expected slowdown than do big jobs.

Prof. Vittoria de Nitto Personè

16


16

Slowdown vs Response time

Response Time tends to be representative of the performance of just **a few** jobs — the bigger ones
tends to emphasize the performance of the really big jobs, since they count the most in the mean, since their response time tends to be the greatest (emphasized for heavy-tail distr.)

Slowdown tends to be representative of the performance of **most** jobs — because it is dominated by the performance of the large number of small jobs.

we would like to make $E(T_S(x))$ smaller for small x



Prof. Vittoria de Nitto Personè

17

17


Processor sharing

we would like to make $E(T_S)$ smaller for small x

How do we do this if we DON'T know job sizes?

two reasons historically why CPU scheduling is (approximately) processor-sharing

1. in a multi-resource system (including a CPU, disk, memory, etc.) it is useful to have many jobs running simultaneously (rather than just one job at a time) because jobs requiring different resources can be overlapped to increase throughput.
2. PS is a good way to get small jobs out fast, given that we don't know the size of the jobs.



Prof. Vittoria de Nitto Personè

18

18

Processor sharing

should be better than FIFO with respect to $E(T_S)$, because PS gets small jobs out faster, and PS should be a lot better than FIFO with respect to $E(sd)$!

$$Pr\{N_S = n\}^{M/G/1/PS} = \rho^n (1 - \rho) = Pr\{N_S = n\}^{M/M/1/FIFO}$$


$$E(N_S)^{M/G/1/PS} = \frac{\rho}{1 - \rho} = E(N_S)^{M/M/1/FIFO}$$

$$E(T_S)^{M/G/1/PS} = \frac{E(S)}{1 - \rho} = E(T_S)^{M/M/1/FIFO}$$

PS is better than FIFO when $C^2 > 1$


the M/G/1/PS queue is insensitive to the variability of the service time distribution, G

2 arrivi simultanei che richiedono 1 s di servizio



$$E(T_S) = (1+2)/2 = 1.5 \text{ s}$$

$E(T_S) ?$



$$E(T_S) = 2 \text{ s}$$

PS può essere peggio su alcune sequenze di job!

Processor sharing

$$E(T_S(x))^{M/G/1/PS} = \frac{x}{1-\rho}$$

$$E(sd(x))^{M/G/1/PS} = \frac{1}{1-\rho}$$

all jobs have same slowdown: PS as “FAIR” scheduling

$$E(sd(x))^{M/G/1/abstract} = 1 + \frac{\frac{\lambda}{2} E(S^2)}{x(1-\rho)}$$

all the preemptive non-size-based scheduling policies produce the same mean slowdown for all job sizes

$$E(sd(x))^{M/G/1/preempt-non-size-based} = \frac{1}{1-\rho}$$

We would like to get lower slowdowns for the smaller jobs

But how can we give preference to the smaller jobs if we don't know job size?

we do know a job's age (CPU used so far), and age is an indication of remaining CPU demand


If the job size distribution has DFR (e.g. Pareto distribution) then the greater the job's age, the greater its expected remaining demand

→ give preference to jobs with low age (younger jobs) and this will have the effect of giving preference to jobs which we expect to be small

(heavy tail: leggere par. 20.7 !)

Scheduling in Unix: Foreground-Background scheduling

Analytical models
scheduling



Prof. Vittoria de Nitto Personè

23