



Performance Modeling of Computer Systems and Networks

Prof. Vittoria de Nitto Personè

Continuous Random Variates: applications

Università degli studi di Roma Tor Vergata
Department of Civil Engineering and Computer Science Engineering

Copyright © Vittoria de Nitto Personè, 2021
<https://creativecommons.org/licenses/by-nc-nd/4.0/>



1

Discrete Simulation
Continuous RV applications

For our application framework, we will look at:

- arrival process model
- service process model


Prof. Vittoria de Nitto Personè

2

2

Arrival process model

- Model *interarrival* times as RV sequence R_1, R_2, R_3, \dots
- Construct corresponding arrival times A_1, A_2, A_3, \dots defined by
 $A_0 = 0$ and $A_i = A_{i-1} + R_i \quad i=1, 2, \dots$
- by induction, $A_i = R_1 + R_2 + \dots + R_i \quad i=1, 2, \dots$
- since $R_i > 0$, $0 = A_0 < A_1 < A_2 < A_3 < \dots$



Prof. Vittoria de Nitto Personè

3

3

Example

- programs `ssq2` and `ssq3` generate job arrivals in this way, where R_1, R_2, R_3, \dots are *Exponential*($1/\lambda$).
 In both programs, the arrival rate is equal to $\lambda = 0.5$ jobs per unit time

```
double GetArrival()
{ static double arrival = START;
  SelectStream(0);
  arrival += Exponential(2.0);
  return (arrival);}
```

Prof. Vittoria de Nitto Personè

4

4

Example

- programs `sis3` and `sis4` generate demand instances in this way, with *Exponential*($1/\lambda$) interdemand times.

The demand rate corresponds to an average of

- $\lambda = 30.00$ actual demands per time interval in `sis3`
- $\lambda = 120.00$ potential demands per time interval in `sis4`

```
double GetDemand(long *amount)
/*
-----
*/
/* generate a demand instance with rate 120
* generate the next demand instance (time) with rate 30 per time
* and generate a corresponding demand amount
* interval and exactly one unit of demand per demand instance
*/
-----
*/
{
    static double time = START;
    static double time = START;

    SelectStream(0);
    SelectStream(0);
    time += Exponential(1.0 / 120.0); /* demand instance */
    SelectStream(2);
    return (time);
    *amount = Geometric(0.2); /* demand amount */
}
return (time);
}
```

Prof. Vittoria de Nitto Personè

5

5

Def stationary arrival process

If R_1, R_2, R_3, \dots is an *iid* sequence of positive interarrival times with $E[R_i] = 1/\lambda > 0$, then the corresponding sequence of arrival times A_1, A_2, A_3, \dots is a *stationary arrival process* with rate λ

- also known as
 - Renewal* processes
 - Homogeneous* arrival processes
- arrival rate λ has units “arrivals per unit time”
 - If average interarrival time is 0.1 minutes,
 - then the arrival rate is 10 arrivals per minute
- stationary arrival processes are “convenient fiction”
 - important theoretically
 - good approximation
 - understand for nonstationary
- If the arrival rate λ varies with time, the arrival process is *nonstationary*

Prof. Vittoria de Nitto Personè

6

6

stationary Poisson arrival process

- As in `ssq2`, `ssq3`, `sis3`, `sis4`, with lack of information it is usually most appropriate to assume that the interarrival times are *Exponential*($1/\lambda$)
- If R_1, R_2, R_3, \dots is an *iid* sequence of *Exponential*($1/\lambda$) interarrival times, the corresponding sequence A_1, A_2, A_3, \dots of arrival times is a stationary *Poisson* arrival process with rate λ . Equivalently, for $i=1, 2, 3, \dots$, the arrival time A_i is an *Erlang*($i, 1/\lambda$) random variable

$$A_i = R_1 + R_2 + \dots + R_i$$

Prof. Vittoria de Nitto Personè

7

7

Algorithm 1

Given $\lambda > 0$ and $t > 0$, this algorithm generates a *realization* of a stationary Poisson arrival process with rate λ over $(0, t)$

```

a0 = 0.0;          /* a convention */
n = 0;
while(an < t) {
    an+1 = an + Exponential(1 / λ);
    n++;
}
return a1, a2, a3, . . . , an;

```

Prof. Vittoria de Nitto Personè

8


8

Random Arrivals

We now demonstrate the interrelation between Uniform, Exponential and Poisson random variables.

Consider:

- $t > 0$, defines a fixed time interval $(0, t)$
- n represents the number of arrivals in the interval $(0, t)$
- $r > 0$, is the length of a small subinterval located at random interior to $(0, t)$



Correspondingly:

- $\lambda = n / t$ is the arrival rate
- $p = r / t$ is the probability that a particular arrival will be in the subinterval
- $np = nr / t = \lambda r$ is the expected number of arrivals in the subinterval

Prof. Vittoria de Nitto Personè

9

9


random arrivals \rightarrow Poisson

Theorem 1

Let:

- A_1, A_2, A_3, \dots be an iid sequence of *Uniform*(0, t) random variables ("unsorted" arrivals).
- the discrete random variable X be the number of A_i that fall in a fixed subinterval of length $r = pt$ interior to $(0, t)$

If n is large and r / t small, X is indistinguishable from a *Poisson*(λr) random variable with $\lambda = n / t$



Prof. Vittoria de Nitto Personè

10

10

Conclusions on random arrivals

- if *many* arrivals occur *at random* with a rate of λ , the number of arrivals X that will occur in an interval of length r is $Poisson(\lambda r)$

- The probability of x arrivals in an interval with length r is

$$Pr(X = x) = \frac{e^{-\lambda r} (\lambda r)^x}{x!} \quad x = 0, 1, 2, \dots$$


- The probability of no arrivals is: $Pr(X = 0) = e^{-\lambda r}$
- The probability of at least one arrival is:

$$Pr(X > 0) = 1 - Pr(X = 0) = 1 - e^{-\lambda r}$$

For a fixed λ , the probability of at least one arrival increases with increasing interval length r

Random Arrivals → Exponential Interarrivals

- If R represents the time between consecutive arrivals, the possible values of R are $r > 0$
- Consider arrival time A_i selected at random and an interval of length r beginning at A_i



- $R = A_{i+1} - A_i$ will be less than r iff there is at least one arrival in this interval
- the cdf of R is

$$Pr(R \leq r) = Pr(\text{at least one arrival in } r) = 1 - e^{-\lambda r}$$
- R is an *Exponential*($1/\lambda$) random variable

Theorem 2

If arrivals occur at random with rate λ , the corresponding interarrival times form an iid sequence of *Exponential*($1/\lambda$) RVs.

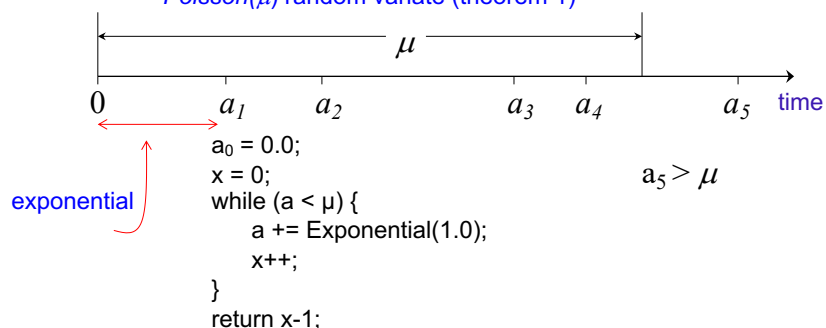
This result justifies the use of *Exponential* interarrival times in programs *ssq2*, *ssq3*, *sis3*, *sis4*

- If we know only that arrivals occur at random with a constant rate λ , the function *GetArrival* in *ssq2* and *ssq3* is appropriate
- If we know only that demand instances occur at random with a constant rate λ , the function *GetDemand* in *sis3* and *sis4* is appropriate

Generating Poisson random variates

Observation:

- If arrivals occur at random with rate $\lambda=1$
- the number of arrivals X in an interval of length μ will be a *Poisson*(μ) random variate (theorem 1)



Summary of Poisson arrival processes

Given a fixed time interval $(0, t)$, there are two ways of generating a realization of a stationary Poisson arrival process with rate λ

1. Generate the number of arrivals: $n = \text{Poisson}(\lambda t)$
Generate a $\text{Uniform}(0, t)$ random variate sample of size n and sort to form $0 < a_1 < a_2 < a_3 < \dots < a_n$
 2. use algorithm 1 with $\text{Exponential}(1/\lambda t)$
- Statistically, the two approaches are equivalent
 - The first approach is computationally more expensive, especially for large n
 - The second approach is always preferred


Prof. Vittoria de Nitto Personè

15

15

Summary of Poisson arrival processes

- The *mode* of the exponential distribution is 0
→ A stationary Poisson arrival process exhibits “clustering”



The stationary Poisson arrival process generalizes to

- a stationary arrival process when exponential interarrival times are replaced by any continuous RV with positive support
- a nonstationary Poisson arrival process when λ varies over time

Prof. Vittoria de Nitto Personè

16

16

For our application framework, we will look at:

- arrival process model
- service process model

17

Service Process Models

differently from the case of arrival processes, there are no well-defined “default”, only application-dependent **guidelines**:

- $Uniform(a, b)$ service times are usually inappropriate since they rarely “cut off” at a maximum value b
- Service times are positive, so they cannot be $Normal(\mu, \sigma)$ unless truncated to positive values
- Positive probability models “with tails”, such as the $Lognormal(a, b)$ distribution, are candidates
 - jobs UNIX
 - web file size
 - Internet topology
 - IP packet flow
 - ...
- If service times are the sum of n iid $Exponential(b)$ sub-task times, then the $Erlang(n, b)$ model is appropriate

18

Program ssq4

- `ssq4` is based on program `ssq3`, but with a more realistic *Erlang*(5, 0.3) service time model
The corresponding service rate is $2/3$
- As in program `ssq3`, `ssq4` uses *Exponential*(2) random variate interarrivals.
The corresponding arrival rate is $1/2$


Prof. Vittoria de Nitto Personè

19

19

Example

- For both `ssq3` and `ssq4`, the arrival rate is $\lambda=0.5$ and the service rate is $\nu=2/3 \approx 0.667$
- The distribution of service times for two programs is very different



first conjecture how the steady-state statistics will differ for these two programs and then investigate the correctness of your conjecture via discrete-event simulation

Prof. Vittoria de Nitto Personè

20

20

Example

- suppose using a $Normal(1.5, 2.0)$ random variable to model service times
- Truncate distribution so that
 - Service times are non-negative ($a=0$)
 - Service times are less than 4 ($b=4$)

```
α = cdfNormal(1.5, 2.0, a);    /* a is 0.0 */
β = 1.0 - cdfNormal(1.5, 2.0, b); /* b is 4.0 */
```

- the result: $\alpha = 0.2266$ and $\beta = 0.1056$

NB

the *truncated* $Normal(1.5, 2.0)$ random variable has a mean of 1.85 (not 1.5) and a standard deviation of 1.07 (not 2.0)

Why is the mean increased?????

Prof. Vittoria de Nitto Personè

21

21

Constrained Inversion

once α and β are determined, the corresponding truncated random variate can be generated by using constrained inversion

```
u = Uniform(α, 1.0 - β);
return F-1(u);
```

Exercise:

- generate n truncated random variates;
- compute sample mean and standard deviation (by Welford)
- verify for which n the statistics converge to the theoretical values

Prof. Vittoria de Nitto Personè

22

22

Example (cont.)

- The idf capability in `rvms` can be used to generate the *truncated Normal*(1.5,2.0) RV

```

 $\alpha = 0.2266274;$ 
 $\beta = 0.1056498;$ 
 $u = \text{Uniform}(\alpha, 1.0 - \beta);$ 
return idfNormal(1.5, 2.0, u);

```




Figure shows $u=0.7090$ and $x=2.601$

Prof. Vittoria de Nitto Personè

23

23

Exercises

- Exercise 7.3.1

Prof. Vittoria de Nitto Personè

24

24